

# Speedup Techniques for Hyperparameter Optimization

## Overview of Multi-Fidelity Optimization

Bernd Bischl   Frank Hutter   Lars Kotthoff  
Marius Lindauer   Joaquin Vanschoren

# Motivating Example

- One possible cheap approximation of an expensive function: use a data subset
  - ▶ Many cheap evaluations on small subsets
  - ▶ Few expensive evaluations on the full data




# Motivating Example

- One possible cheap approximation of an expensive function: use a data subset

- ▶ Many cheap evaluations on small subsets
- ▶ Few expensive evaluations on the full data



- E.g.: Support Vector Machines (SVM) on MNIST dataset (hyperparameters:  $C$ ,  $\gamma$ )




# Motivating Example

- One possible cheap approximation of an expensive function: use a data subset

- ▶ Many cheap evaluations on small subsets
- ▶ Few expensive evaluations on the full data



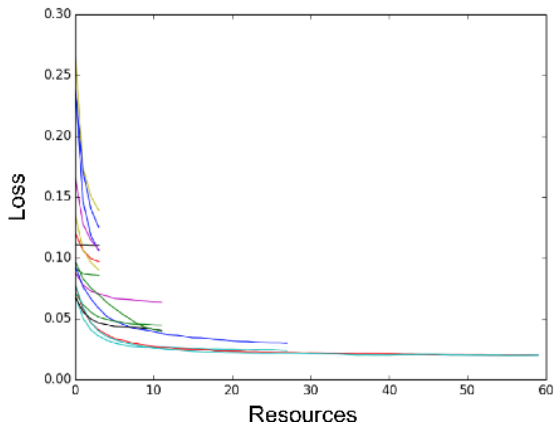
- E.g.: Support Vector Machines (SVM) on MNIST dataset (hyperparameters:  $C$ ,  $\gamma$ )



→ up to 1000x speedups over blackbox optimization on full data [Klein et al, AISTATS 2017]

## Motivating Example 2: Shorter Runs of Anytime Algorithms


- Performance with shorter runs of an anytime algorithm (such as SGD):



# Multi-Fidelity Optimization In General

Exploit cheap approximations of an expensive blackbox function  $\rightarrow$  afford more configurations


- Idea: eliminate poor configurations early, allocate more resources to promising ones.



# Multi-Fidelity Optimization In General

Exploit cheap approximations of an expensive blackbox function  $\rightarrow$  afford more configurations


- Idea: eliminate poor configurations early, allocate more resources to promising ones.
- Possible Resources:
  - ▶ Data subset size
  - ▶ Runtime / # epochs / # iterations



# Multi-Fidelity Optimization In General

Exploit cheap approximations of an expensive blackbox function  $\rightarrow$  afford more configurations

- Idea: eliminate poor configurations early, allocate more resources to promising ones.
- Possible Resources:
  - ▶ Data subset size
  - ▶ Runtime / # epochs / # iterations
  - ▶ Downsampled size of images in object recognition
  - ▶ Depth / width of neural networks






# Multi-Fidelity Optimization In General

Exploit cheap approximations of an expensive blackbox function  $\rightarrow$  afford more configurations


- Idea: eliminate poor configurations early, allocate more resources to promising ones.
- Possible Resources:
  - ▶ Data subset size
  - ▶ Runtime / # epochs / # iterations
  - ▶ Downsampled size of images in object recognition
  - ▶ Depth / width of neural networks
  - ▶ Number of trees
  - ▶ Number of features
  - ▶ Number of cross validation folds



# Multi-Fidelity Optimization In General

Exploit cheap approximations of an expensive blackbox function → afford more configurations

- Idea: eliminate poor configurations early, allocate more resources to promising ones.
- Possible Resources:
  - ▶ Data subset size
  - ▶ Runtime / # epochs / # iterations
  - ▶ Downsampled size of images in object recognition
  - ▶ Depth / width of neural networks
  - ▶ Number of trees
  - ▶ Number of features
  - ▶ Number of cross validation folds
- ▶ General concept, applicable even in fields outside ML, e.g., fluid simulation:
  - ★ Number of particles
  - ★ Time scale of simulation



# General Remarks on Multi-Fidelity Optimization

- Often, we have a **choice** which resources we use as budget

# General Remarks on Multi-Fidelity Optimization

- Often, we have a **choice** which resources we use as budget
- For multi-fidelity optimization to be helpful, **performance with low budgets should be informative about performance with high budgets**

# General Remarks on Multi-Fidelity Optimization

- Often, we have a **choice** which resources we use as budget
- For multi-fidelity optimization to be helpful, **performance with low budgets should be informative about performance with high budgets**
- In the simplest case: good with low resources  $\leftrightarrow$  good with high resources.
  - ▶ In practice, this is of course not always true

# How Useful is the Cheap Approximation? The Rank Correlation

Given:

- A set of configurations  $\Lambda = \{\lambda_1, \dots, \lambda_n\}$
- The performances  $f(\lambda_1), \dots, f(\lambda_n)$  on the expensive black box
- The performances  $g(\lambda_1), \dots, g(\lambda_n)$  on a cheap approximation of the black box

# How Useful is the Cheap Approximation? The Rank Correlation

Given:

- A set of configurations  $\Lambda = \{\lambda_1, \dots, \lambda_n\}$
- The performances  $f(\lambda_1), \dots, f(\lambda_n)$  on the expensive black box
- The performances  $g(\lambda_1), \dots, g(\lambda_n)$  on a cheap approximation of the black box

We compute the **Spearman rank correlation** between  $[f(\lambda_1), \dots, f(\lambda_n)]$  and  $[g(\lambda_1), \dots, g(\lambda_n)]$

- If this is high (in the extreme: 1), the relative ranking of the configurations is the same on  $f$  and  $g$ 
  - ▶ In that case, we can optimize cheaply on  $g$  and also obtain an optimum for  $f$
- If it is low ( $\approx 0$ ), optimizing  $g$  does not tell us anything about  $f$

# How Useful is the Cheap Approximation? The Rank Correlation

Given:

- A set of configurations  $\Lambda = \{\lambda_1, \dots, \lambda_n\}$
- The performances  $f(\lambda_1), \dots, f(\lambda_n)$  on the expensive black box
- The performances  $g(\lambda_1), \dots, g(\lambda_n)$  on a cheap approximation of the black box

We compute the **Spearman rank correlation** between  $[f(\lambda_1), \dots, f(\lambda_n)]$  and  $[g(\lambda_1), \dots, g(\lambda_n)]$

- If this is high (in the extreme: 1), the relative ranking of the configurations is the same on  $f$  and  $g$ 
  - ▶ In that case, we can optimize cheaply on  $g$  and also obtain an optimum for  $f$
- If it is low ( $\approx 0$ ), optimizing  $g$  does not tell us anything about  $f$

Goal: find approximations  $g$  that are very cheap but have high rank correlations with  $f$



## Questions to Answer for Yourself / Discuss with Friends

- **Repetition.** Which cheap approximation is better in this hypothetical case?
  - ▶ Downscaling images (5x cheaper, rank correlation of 0.8)
  - ▶ Less epoch of SGD (4x cheaper, rank correlation of 0.75)
- **Discussion.** Can you think of an application of your interest where you would likely have a good multi-fidelity approximation?